




**ISC**

High Performance

REINVENTING

HPC

MAY 12 – 16, 2024 | HAMBURG, GERMANY





# EAR ISC2024 tutorial: Advanced use cases and optimization

Julita Corbalan ([julita.corbalan@eas4dc.com](mailto:julita.corbalan@eas4dc.com))



Benjamin Czaja ([benjamin.czaja@surf.nl](mailto:benjamin.czaja@surf.nl))





# Advance use cases and optimization

- Codes in /projects/0/energy-course and GIT
- GIT: <https://github.com/sara-nl/ISC-2024-EAR-tutorial/tree/main>
- Get the examples and test them
  - [https://github.com/sara-nl/ISC-2024-EAR-tutorial/tree/main/tutorials/monitoring\\_ear](https://github.com/sara-nl/ISC-2024-EAR-tutorial/tree/main/tutorials/monitoring_ear)
  - GROMACS singularity
  - PyTorch
  - Palabos
- Data visualization
  - <https://github.com/sara-nl/ISC-2024-EAR-tutorial/tree/main/tutorials/visualization>
  - Grafana (using local installation)
  - ear-job-analytics tool

# Application characterization/optimization



# Application characterization/optimization





# Use cases

- Singularity
  - GROMACS (monitoring)
- CPU apps
  - Palabos: monitoring and dynamic optimization
  - NPB: dynamic optimization
- GPU apps
  - GROMACS singularity (monitoring)
  - PyTorch ( min\_energy)
- All of them, understand, visualize and compare energy efficiency when doing dynamic optimization





# GROMACS-GPU: Singularity + EAR

- Singularity/Apptainer
  - <https://apptainer.org/>
  - <https://docs.sylabs.io/guides/3.5/user-guide/introduction.html>
  - [https://catalog.ngc.nvidia.com/orgs/hpc/collections/nvidia\\_hpc/entities](https://catalog.ngc.nvidia.com/orgs/hpc/collections/nvidia_hpc/entities)
- Singularity containers allow applications to use host services (such as ear)
  - Paths must be binded
  - Environment variables must be defined
  - [https://github.com/sara-nl/ISC-2024-EAR-tutorial/blob/main/tutorials/monitoring\\_ear/GROMACS\\_SINGULARITY\\_GPU.sh](https://github.com/sara-nl/ISC-2024-EAR-tutorial/blob/main/tutorials/monitoring_ear/GROMACS_SINGULARITY_GPU.sh)




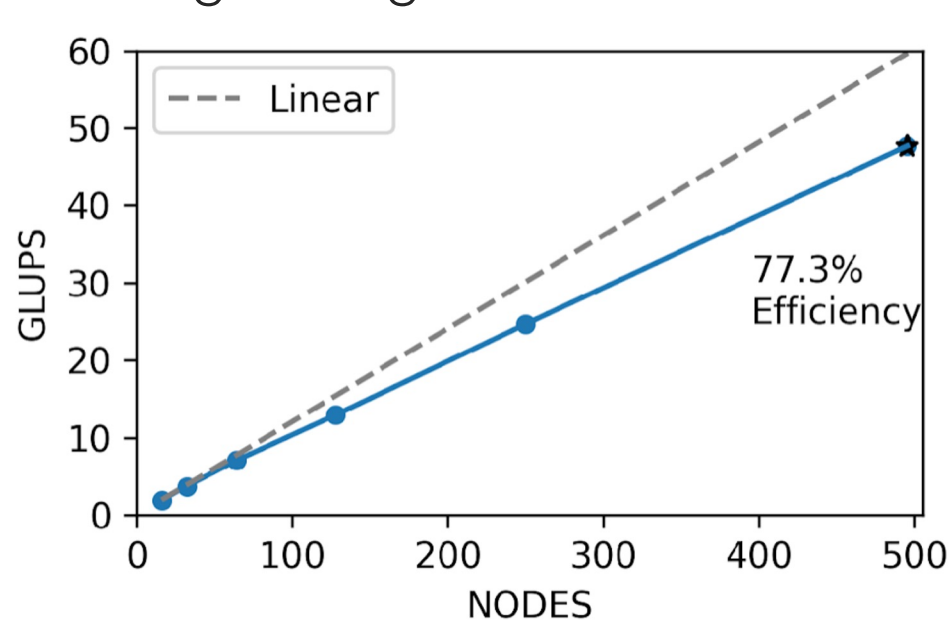
# optimization

- Codes in /projects/0/energy-course
- GIT: <https://github.com/sara-nl/ISC-2024-EAR-tutorial/tree/main>
- Get the examples, add EAR **monitoring and min\_energy** policy and compare
  - `-ear=on` → monitoring
  - `-ear-policy=min_energy` → selects min\_energy policy
  - Run 2 steps in the same job to guarantee both runs are executed in the same node (s)
  - [https://github.com/sara-nl/ISC-2024-EAR-tutorial/tree/main/tutorials/monitoring\\_ear](https://github.com/sara-nl/ISC-2024-EAR-tutorial/tree/main/tutorials/monitoring_ear)
  - NPB : Rome vs Genoa
  - Palabos (use 1 and 4 nodes, node input\_1\_node\_XL.xml input file)
    - Rome vs Genoa
  - PyTorch
    - Get some extra savings when running in exclusive mode
    - `export EAR_JOB_EXCLUSIVE_MODE=1`



# Palabos: Lattice-Boltzmann Solver

## Strong Scaling Benchmark




- An average electric car consumes about 0.2 kwh/km
- 495 node case (which ran for 9 minutes) just drove a car 150 km!

# Palabos:




## Strong Scaling Benchmark

- Per-node, Per-iteration "traces"
- Node Power
- Avg CPU Freq (node)
- Main memory BW (GB/s)
- CPI (Cycles per instruction)
- MPI% (percentage spent in MPI calls)
- I/O (network communication)

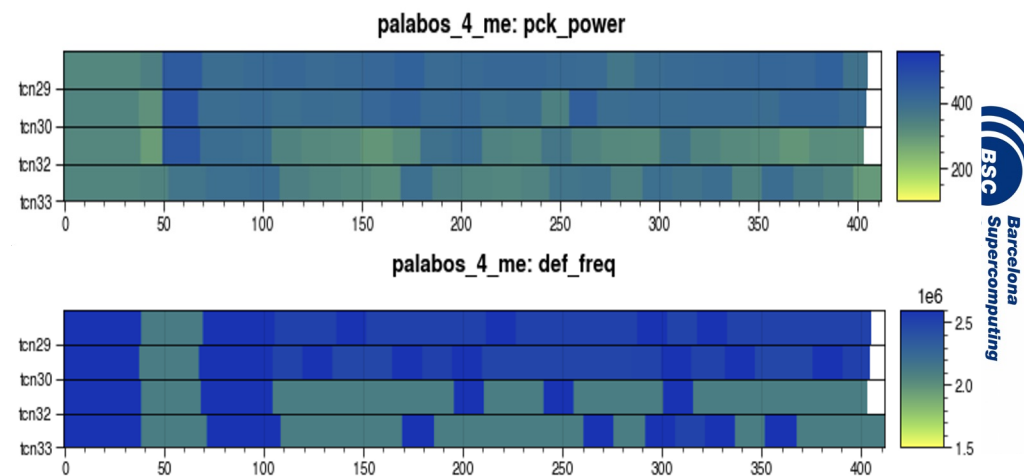


# Data visualization

- <https://github.com/sara-nl/ISC-2024-EAR-tutorial/blob/main/tutorials/visualization/README.md>
- ear-job-analytics
  - Requires loops in DB: export EARL\_REPORT\_LOOPS=1
  - Use ear-job-anaytics directly or use create\_traces.sh script
- Grafana
  - Running al DC and executing SQL queries (more powerful, but depends on DC)
  - Local installation:
    - Grafana server installed and running locally
    - Data gathered in CSV format with eacct and using CSV plugin
    - Not mandatory but more information if loops are in DB
    - Can be use also without EAR DB: -ear-user-db=filename



Application metrics



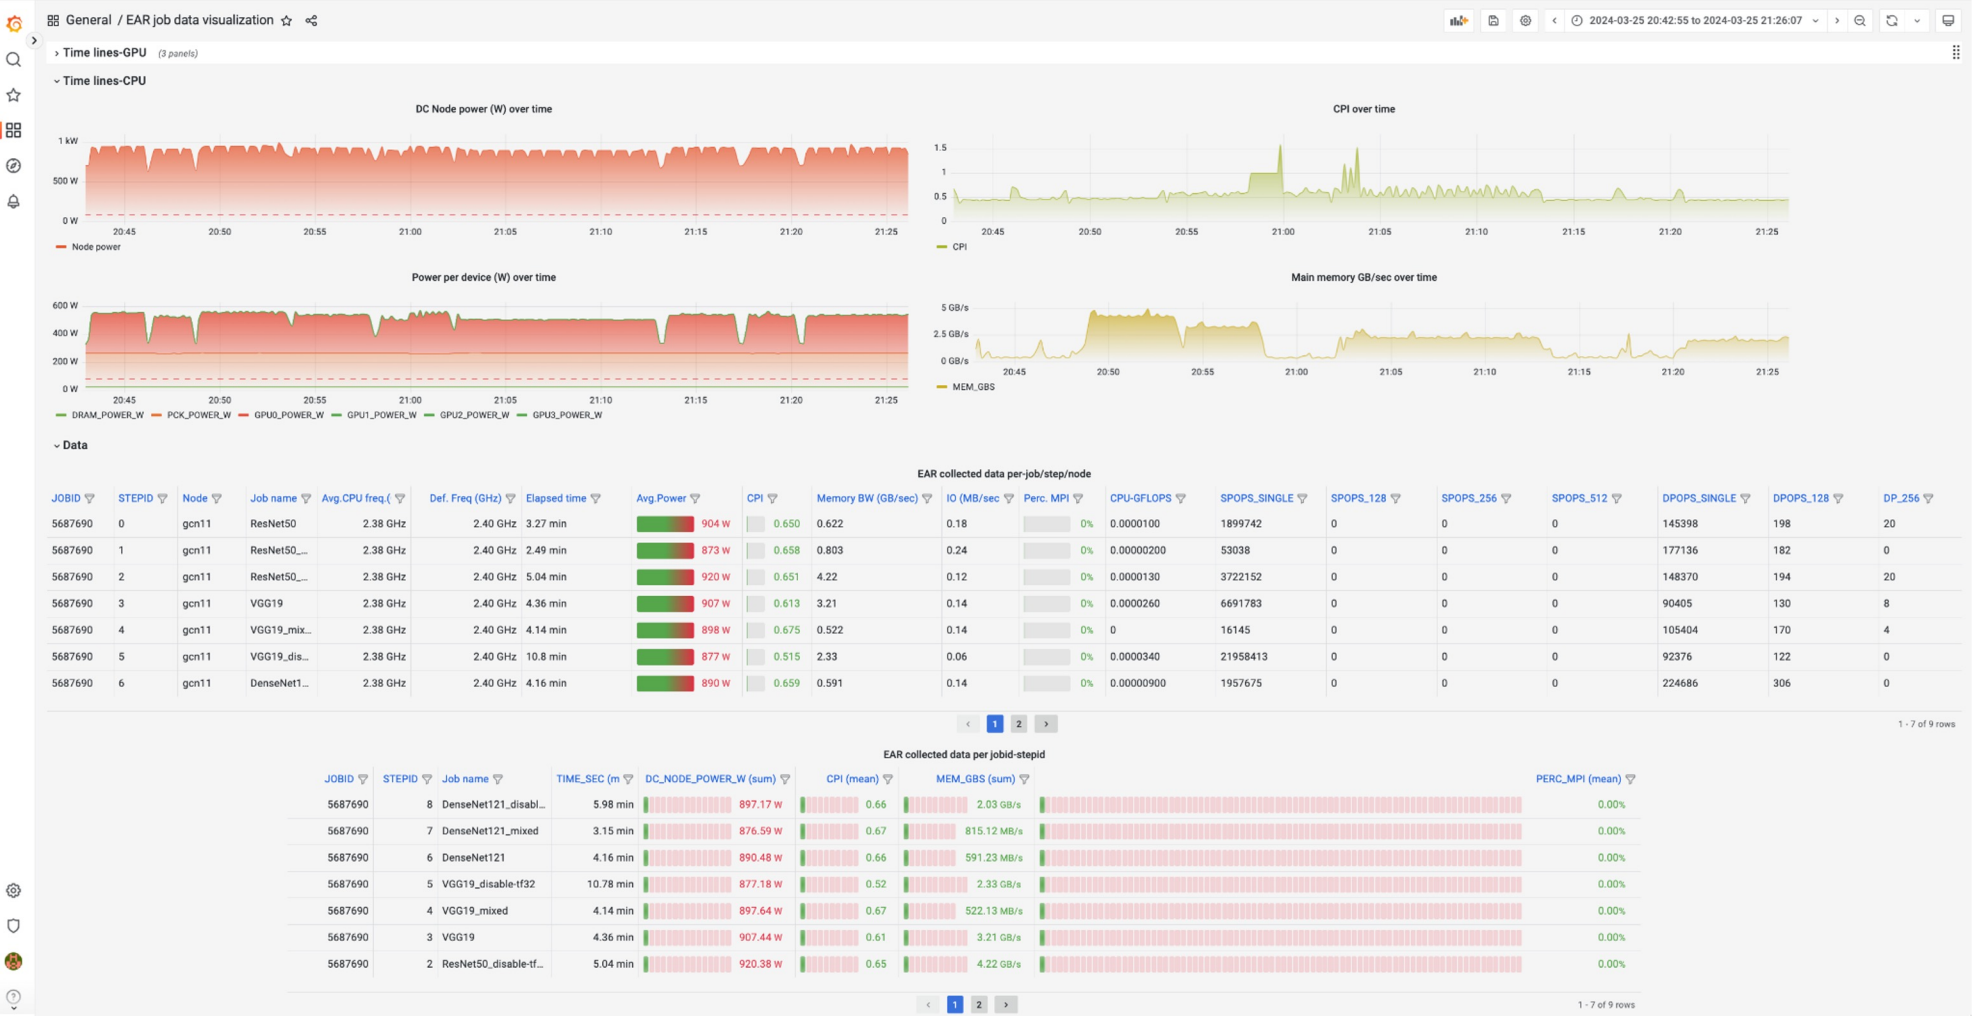
Runtime optimization



# Visualization with Grafana

- Steps to use EAR data are in grafana with CSV are
  - Have a local grafana installation with CSV plugins supported
  - Export EAR application data in csv format using eacct (-l -c option) or `-ear-user-db` flag
  - Export EAR application runtime data in csv format using eacct (-r -c option) or `-ear-user-db` flag
  - Add a source data based on a local file (Public folder)
    - `julita.corbalan$ cp tensorflow.csv ear_data_apps.csv`
    - `julita.corbalan$ cp tensorflow_loops.csv ear_data_loops.csv`
  - Import the EAR json file with the dashboards for data visualization
    - "EAR job data visualization.json"
  - Reload the dashboards





Go to ISC tutorial github and do the exercises

<https://github.com/sara-nl/ISC-2024-EAR-tutorial/tree/main/tutorials/policies>



ISC 2024 | MAY 12 – 16 | HAMBURG, GERMANY | #ISC24